



Grant Agreement: 287829

Comprehensive Modelling for Advanced Systems of Systems

C O M P A S S

## CML Interpreter Design Document

Technical Note Number: DXX

Version: 0.3

Date: Month Year

Public Document

<http://www.compass-research.eu>

<sup>11</sup> **Contributors:**

<sup>12</sup> Anders Kaels Malmos, AU

<sup>13</sup> **Editors:**

<sup>14</sup> Peter Gorm Larsen, AU

<sup>15</sup> **Reviewers:**

## 16 Document History

Ver	Date	Author	Description
0.1	25-04-2013	Anders Kaels Malmos	Initial document version
0.2	06-03-2014	Anders Kaels Malmos	Added introduction and domain description
0.3	26-03-2014	Anders Kaels Malmos	Added draft of static and dynamic structure for the core interpreter

## 18 **Abstract**

19 This document describes the overall design of the CML interpreter and provides an  
20 overview of the code structure targeting developers. It assume a basic knowledge of  
21 CML.

## Contents

23	<b>1 Introduction</b>	<b>6</b>
24	1.1 Problem Domain . . . . .	6
25	1.2 Definitions . . . . .	8
26	<b>2 Software Layers</b>	<b>9</b>
27	<b>3 Layer design and Implementation</b>	<b>9</b>
28	3.1 The Core Layer . . . . .	10
29	3.2 The Dynamic Model . . . . .	16
30	3.3 The IDE Layer . . . . .	18

# 1 Introduction

This document is targeted at developers and describes the overall design of the CML simulator, it is not a detailed description of every part of the source code. This kind of documentation is done in Javadoc and can be generated automatically from the code. It is assumed that common design patterns are known like [GHJV94] and a basic understanding of CML.

## 1.1 Problem Domain

The goal of the interpreter is to enable simulation/animation of a given CML model and be able to visualize this in the Eclipse IDE Debugger. CML has a UTP semantics defined in [BGW13] which dictates the interpretation. Therefore, the overall goal of the CML interpreter is to adhere to the semantic rules defined in those documents and to visualize this in the Eclipse Debugger.

In order to get a high level understanding of how CML is interpreted without knowing all the details, a short illustration of how the interpreter represents and evolves a CML model is given below.

In Listing 1 a CML model consisting of three CML processes is given. It has a R (Reader) process which reads a value from the inp channel and writes it on the out channel. The W (Writer) process writes the value 1 to the inp channel and finishes. The S (System) process is a parallel composition of these two processes where they must synchronize all events on the inp channel.

```

51 channels
52 inp : int
53 out : int
54
55 process W =
56 begin
57   @ inp!1 -> Skip
58 end
59
60 process R =
61 begin
62   @ inp?x -> out!x -> Skip
63 end
64
65 process S = W [|{$inp$}] R

```

Listing 1: A process S composed of a parallel composition of a reader and writer process

The interpretation of a CML model is done through a series of steps/transitions starting from a given entry point. In figure 1 the first step in the interpretation of the model is shown, it is assumed that the S process is given as the starting point. Processes are represented as a circle along with its current position in the model. Each step of the execution is split up in two phases, the inspection phase and the execution phase. The dashed lines represent the environment (another actor that invokes the operation e.g. a human user or another process) initiating the phase.




Figure 1: Initial step of Listing 1 with process S as entry point.

The inspection phase determines the possible transitions that are available in the next step of execution. The result of the inspection is shown as a set of transitions below “inspect”. As seen on figure Figure 1 process P starts out by pointing to the parallel composition constructs, this construct has a semantic begin rule which does the initialization needed. In the figure Figure 1 that rule is named  $\tau(\text{par-begin})$  and is therefore returned from the inspection. The reason for the name  $\tau(\dots)$  is that transitions can be either observable or silent, so in principle any  $\tau$  transition is not observable from the outside of the process. However, in the interpreter all transitions flows out of the inspection phase. When the inspection phase has completed, the execution phase begins. The execution phase executes one of the transitions returned from the inspection phase. In this case, only a single transition is available so the  $\tau(\text{par-begin})$  is executed which creates the two child processes. The result of each of the shown steps are the first configuration shown in the next step. So in this case the resulting process configuration of Figure 1 is shown in figure Figure 2.

The second step on Figure 2 has a more interesting inspection phase. According to the parallel composition rule, we have that any event on the  $\text{inp}$  channel must be synchronized, meaning that W and R must only perform transition that involves  $\text{inp}$  channel events synchronously.

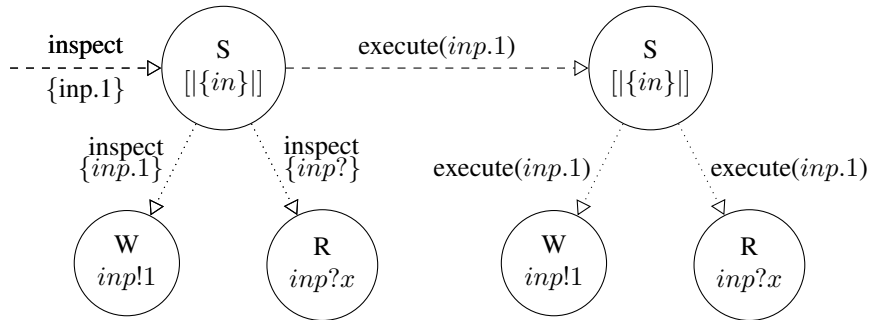


Figure 2: Second step of Listing 1 with S as entry point.

Therefore, when P is inspected it must inspect its child processes to determine the possible transitions. In this case W can perform the  $\text{inp.1}$  event and R can perform any event on  $\text{inp}$  and therefore, the only possible transition is the one that performs the  $\text{inp.1}$  event. This is then given to the execution phase which result in the  $\text{inp.1}$  event and moves both child processes into their next state.

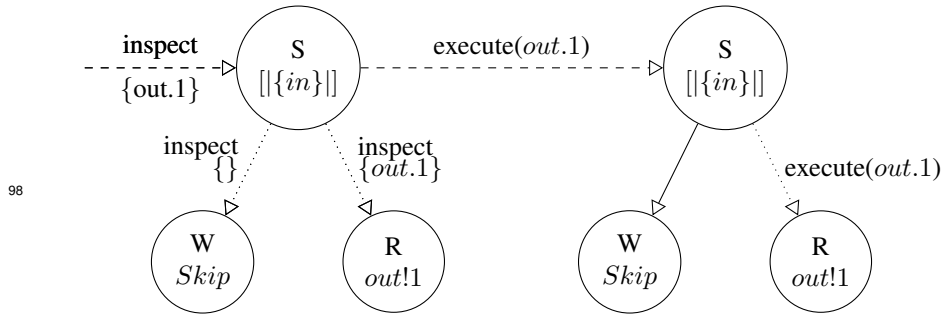


Figure 3: Third step of Listing 1 with S as entry point

In the third step on figure Figure 3 W is now Skip which means that it is successfully terminated. The inspection for W therefore results in an empty set of possible transitions. R is now waiting for the *out.1* event after 1 was writing to *x* in the last step and therefore returns this transition. The execution phase is a little different and S now knows only to execute R.




Figure 4: Final step of Listing 1 where the parallel composition collapses unto a Skip process

The fourth and final step shown in Figure 4 of the interpretation starts out with both W and R as Skip, this triggers the parallel end rules, which evolves into Skip. S therefore returns the silent transition the triggers this end rule.

## 1.2 Definitions

**Animation** Animation is when the user are involved in taking the decisions when interpreting the CML model

**CML** Compass Modelling Language

**UTP** Unified Theory of Programming (a semantic framework)

**Simulation** Simulation is when the interpreter runs without any form of user interaction other than starting and stoppping.

**trace** A sequence of observable events performed by a behavior.



## 2 Software Layers

This section describes the layers of the CML interpreter. As depicted in figure 5 two highlevel layers exists.




Figure 5: The layers of the CML Interpreter

Each of these components will be described in further detail in the following sections. The major reason behind this layering is that the implementation of the semantics should be independent of the view showing the results.

**Core Layer** This layer has the overall responsibility of interpreting a CML model as described in the operational semantics that are defined in [BGW13] and is located in the package *eu.compassresearch.core.interpreter*

**IDE Layer** Has the overall responsibility of visualizing the outputs of a running interpretation a CML model in the Eclipse Debugger. It is located in the *eu.compassresearch.ide.plugins.interpreter* package. The IDE part is integrating the interpreter into Eclipse, enabling CML models to be debugged through the Eclipse debugger.

## 3 Layer design and Implementation

This section describes the static and dynamic structure of the components involved in interpreting a CML model.

## 3.1 The Core Layer

The core layer is responsible for the overall interpretation of a given CML model. To understand some of the choice made, the design philosophy needs a short word. The design philosophy of the top-level structure is to encapsulate all the classes and interfaces (hence make elements package accessible only when appropriate) that makes up the implementation of the core functionality and only expose those that are needed to utilize the interpreter. This provides a clean separation between the implementation and interface and makes it clear for both the users, which not necessarily wants to know about the implementation details, and developers which parts they need to work with.

In the following section both the static and dynamic model will be described in more details.

### 3.1.1 The Static Model

#### Packages

The following packages defines the top level structure of the core:

**eu.compassresearch.core.interpreter** This package contains all the internal classes and interfaces that defines the core functionality of the interpreter. There is one important public class in the package, namely the **VanillaInterpreterFactory** factory class, that any user of the interpreter must invoke to use the interpreter. This can creates instances of the **CmlInterpreter** interface. Furthermore, this package is split into two seperate source folders, each representing a different logical component. The following folders are present:

**src/main/java** This folder contains all public classes and interfaces as described above.

**src/main/behavior** This folder contains all the internal classes and interfaces that the default interpreter implementation is comprised of. This will be described in more details in Subsection 3.1.1.

**eu.compassresearch.core.interpreter.api** This package and sub-packages contains all the public classes and interfaces that defines the API of the interpreter. Some of the most important entities of this package includes the main interpreter interface **CmlInterpreter** along with the **CmlBehaviour** interface that represents a CML process or action. It corresponds to the circles in the figures of Subsection 1.1.

**eu.compassresearch.core.interpreter.api.events** This package contains all the public components that enable users of the interpreter to subscribe to multiple events (this it not CML channel events) from both **CmlInterpreter** and **CmlBehaviour** instances.

**eu.compassresearch.core.interpreter.api.transitions** This package contains all the possible types of transitions that a **CmlBehaviour** instance can make. This will be explained in more detail in section 3.1.1.

171 **eu.compassresearch.core.interpreter.api.values** This package contains all the val-  
 172 ues used by the CML interpreter. They represent the values of variables and  
 173 constants in a context.

174 **eu.compassresearch.core.interpreter.cosim** Has the responsibility of running a co-  
 175 simulation. A co-simulation can be either between multiple instances of the  
 176 CML interpreter co-simulating a CML model, or a CML interpreter instance co-  
 177 simulating a CML model with a real live system.

178 **eu.compassresearch.core.interpreter.remote** This has the responsibility of exposing  
 179 the CML interpreter to be remote controlled.

180 **eu.compassresearch.core.interpreter.debug** Has the responsibility of controlling a  
 181 debugging sessions, which only includes the Eclipse debugger at this point.

182 **eu.compassresearch.core.interpreter.utility** The utility packages contains reusable  
 183 classes and interfaces that are use across packages.

## 184 The Top Level Elements

185 The top level interfaces and classes of the interpreter structure is depicted in Figure 6,  
 186 followed by a short description of each the depicted components.

187 Before going into details with each element on figure 6 a few things needs mentioning.  
 188 First of all, any CML model has a top level Process. Because of this, the interpreter  
 189 need only to interact with the top level CmlBehaviour instance. This explains the one-  
 190 to-one correspondence between the CmlInterpreter and the CMLBehaviour. However,  
 191 the behavior of top level CmlBehaviour is determined by the binary tree of CmlBe-  
 192 haviour instances that itself and it's child behaviours defines. So in effect, the CmlIn-  
 193 terpreter along with the selection strategy controls every observable transition that any  
 194 CmlBehaviour makes.

195 **CmlInterpreter** The interface exposing the functionality of the interpreter compo-  
 196 nent. This interface has the overall responsibility of interpreting. It exposes  
 197 methods to inspect and execute and it is implemented by the **VanillaCmlInter-**  
 198 **preter** class in the default simulation settings.

199 **CmlBehavior** Interface that represents a behavior specified by either a CML process  
 200 or action. Most importantly it exposes the two methods: *inspect* which calcu-  
 201 lates the immediate set of possible transitions that it currently allows and *execute*  
 202 which takes one of the possible transitions determined by it's supervisor. This  
 203 process is described in Subsection 1.1 where a CmlBehavior is represented as a  
 204 circle in the figures. As seen both in Subsection 1.1 and Figure 6 associations  
 205 between CmlBehavior instances are structured as a binary tree, where a parent  
 206 supervises its child behaviors. In this context supervises means that they control  
 207 the flow of possible transitions and determines when to execute them. The reason  
 208 for this is that is corresponds nicely to the structure of the CML semantics.

209 **SelectionStrategy** This interface has the responsibility of choosing a CmlTransition  
 210 from a given CmlTransitionSet. This could be seen as the last chain in the super-  
 211 visor hierarchy, since this is where all the possible transitions flows to and the  
 212 decision of which one to execute next is taken here. The purpose of this interface

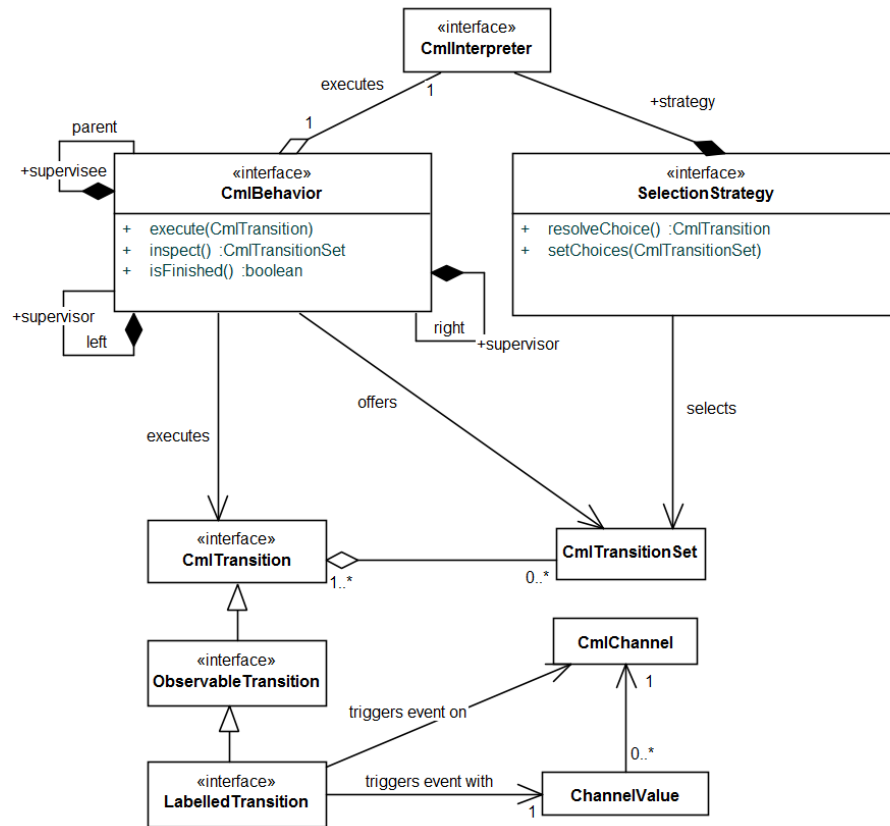


Figure 6: The high level classes and interfaces of the interpreter core component

is to allow different kinds of strategies for choosing the next transition. e.g there is a strategy that picks one at random and another that enables a user to pick.

**CmlTransition** Interface that represents any kind of transition that a CmlBehavior can make. They are not all depicted here and will be described in greater details in ???. But overall, only transitions that implements the ObservableTransition interface can produce an observable trace of a behavior.

**CmlTransitionSet** This is an immutable set of CmlTransition objects and is the return value of the inspect method on a CmlBehavior. The reason for it being immutable is to ensure that calculations never change the input sets.

## The Transitions Model

As described in the previous sections a CML model is represented by a binary tree of CmlBehaviour instances and each of these has a set of possible transitions that they can make. A class diagram of all the classes and interfaces that makes up transitions are shown in figure ??, followed by a description of each of the elements.

A transition taken by a CmlBehavior is represented by a CmlTransition. This represent

228 a possible next step in the model which can be either observable or silent (also called a  
229 tau transition).

230 An observable transition represents either that time passes or that a communication/syn-  
231 chronization event takes place on a given channel. All of these transitions are captured  
232 in the ObservableTransition interface. A silent transitions is captured by the TauTran-  
233 sition and HiddenTransition class and can respectively marks the occurrence of a an  
234 internal transition of a behavior or a hidden channel transition.

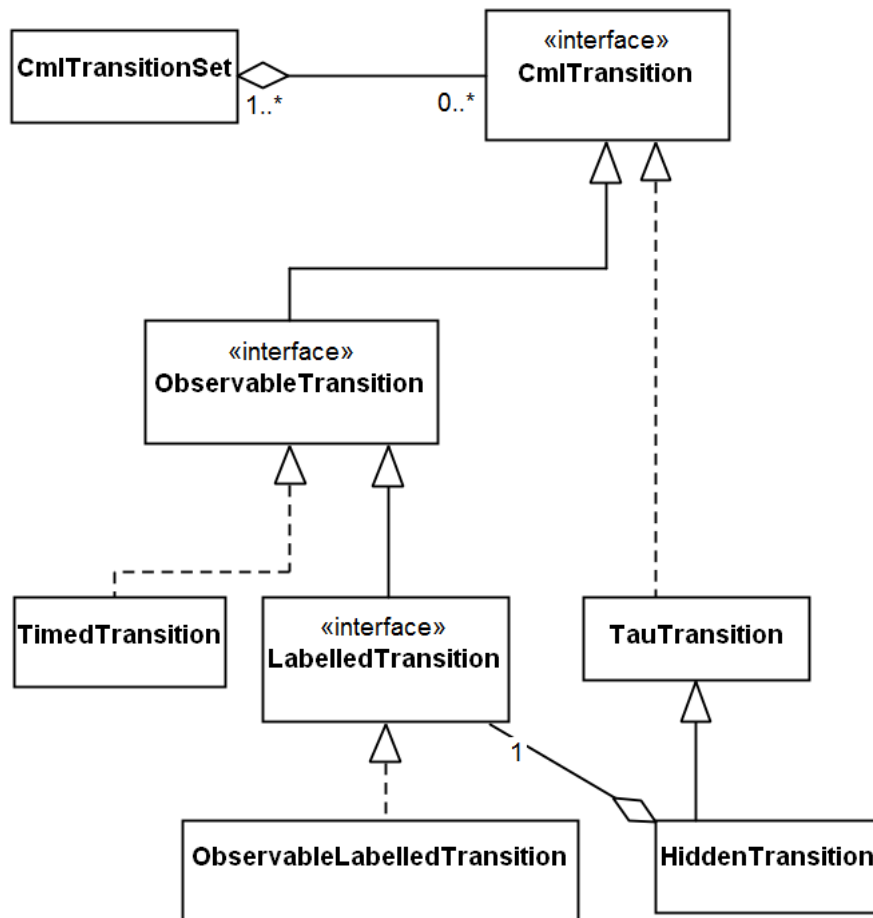


Figure 7: The classes and interfaces that defines transitions

235 **CmlTransition** Represents any possible transition.

236 **CmlTransitionSet** Represents a set of CmlTransition objects.

237 **ObservableTransition** This represents any observable transition.

238 **LabelledTransition** This represents any transition that results in a observable channel  
239 event

240 **TimedTransition** This represents a tock event marking the passage of a time unit.

241 **ObservableLabelledTransition** This represents the occurrence of a observable chan-  
 242 nel event which can be either a communication event or a synchronization event.

243 **TauTransition** This represents any non-observable transitions that can be taken in a  
 244 behavior.

245 **HiddenEvent** This represents the occurrence of a hidden channel event in the form of  
 246 a tau transition.

### 247 The Default CmlBehavior Implementation

248 Actions and processes are both represented by the CmlBehaviour interface. A class dia-  
 gram of the important classes that implements this interface is shown in Figure 8 When

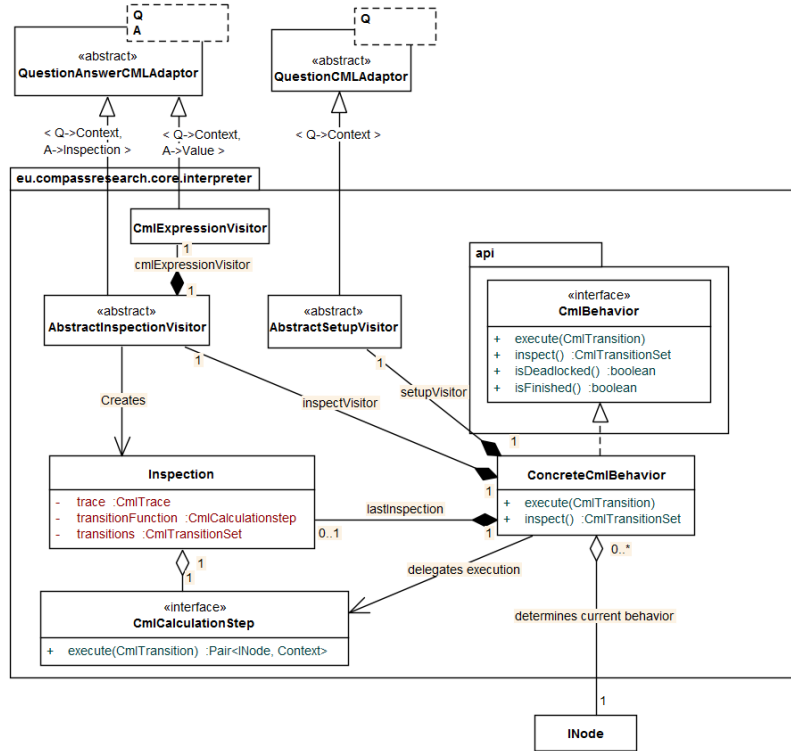


Figure 8: The classes and interfaces making up the default implementation the CmlBehavior interface

249 the interpreter runs in the default operation mode, meaning where only a single inter-  
 250 preter instance runs (opposed to the co-simulation modes where multiple instances of  
 251 the interpreter might run or connected to an externally running system). Then all Cml-  
 252 Behavior instances will be in the form of the ConcreteCmlBehavior class. As described  
 253 above a CmlBehavior has the responsibility to behave as a given action or process.  
 254 However, as shown in Figure 8 the ConcreteCmlBehavior class delegates a large part  
 255 of its responsibility to other classes. The actual behavior of a ConcreteCmlBehavior  
 256 instance is decided by its current INode instance, so when a ConcreteCmlBehavior in-  
 257 stance is created a INode instance must be given. The INode interface is implemented  
 258

by all the CML AST nodes and can therefore be any CML process or action. The actual implementation of the behavior of any process/action is delegated to internal visitor classes as depicted in Figure 8. The used visitors are all extending generated abstract visitors that have the infrastructure to visit any CML AST node. The reason for this structure is to be able to utilize the already generated visitors by the AST-creator (located at <https://github.com/overturetool/astcreator>) that enables traversing of CML AST's.

Here a brief description of each new element depicted in Figure 8:

**CmlExpressionVisitor** This has the responsibility to evaluate CML expressions given a Context.

**AbstractSetupVisitor** This has the responsibility of performing any required setup for a behavior. This visitor is invoked whenever a new INode instance is loaded.

**AbstractAlphabetVisitor** This has the responsibility of creating an Inspection object given the current state of the behavior, which is represented by a INode and a Context object.

**Inspection** Contains the next possible transitions (in a CmlTransitionSet) along with a transition function in the form of a CmlCalculationStep.

**CmlCalculationStep** Responsible for executing the actual behavior that occurs in a transition from one state to another. This is where the actual implementation of the semantics is.

## The Visitors

In figure 9 a more detailed look at the inspection visitor structure is given.

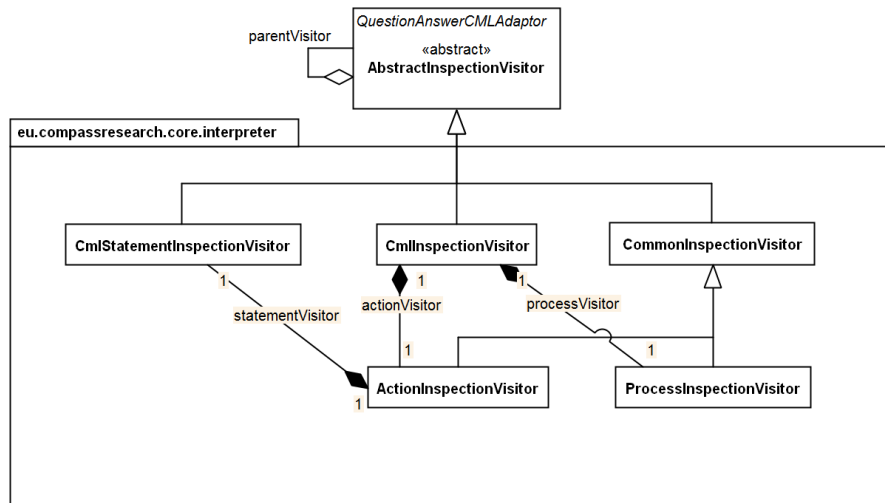


Figure 9: Visitor structure

281 As depicted the visitors are split into several visitors that handle different parts of the  
282 CML language. The sole reason for doing this is to avoid having one large visitor  
283 that handles all the cases. At run-time the visitors are setup in a tree structure. For  
284 the inspection the top most visitor is the CmlInspectionVisitor which then delegates  
285 to either ActionInspectionVisitor or ProcessEvaluationVisitor depending on the given  
286 INode. This structures resembles the structure of the setup visitors.

287 The CmlExpressionVisitor is however a little different from the others. It takes care  
288 of all CML expressions, but delegates the entire subset of VDM expression constructs  
289 that are contained in CML to the DelegateExpressionEvaluator overture class, which  
290 can evaluate VDM expressions. The reason for doing this is of course reuse.

## 291 3.2 The Dynamic Model

292 This section will describe the high-level dynamic model. First of all, in the default oper-  
293 ation mode (as mentioned above a single running instance of the interpreter) the entire  
294 CML interpreter runs in a single thread. This is mainly due to the inherent complex-  
295 ity of concurrent programming. You could argue that since a large part of COMPASS  
296 is about modelling complex concurrent systems, we also need a concurrent interpre-  
297 tation of the models. However, the semantics is perfectly implementable in a single  
298 thread which makes a multi-threaded interpreter optional. There are of course benefits  
299 to a multi-threaded interpreter, but for matters such as the testing and deterministic be-  
300 haviour a single threaded interpreter is much easier to handle and comprehend.

### 301 The Top Execution Loop

302 To start a simulation/animation of a CML model, you first of all need an instance of  
303 the CmlInterpreter interface. This is created through the VanillaInterpreterFactory by  
304 invoking the newInterpreter method with a typechecked AST of the CML model. The  
305 default returned instance is the VanillaCmlInterpreter class. Once a CmlInterpreter is  
306 instantiated the interpretation of the CML model is started by invoking the execute  
307 method.

308 In figure 10 a sequence diagram of the execute method on the VanillaCmlInterpreter  
309 class is depicted.

310 As seen in the figure the execution continues until the top level process is either suc-  
311 cessfully terminated or deadlocked. Each round taken in this loop is one step taken in  
312 the model, where the meaning of a step is explained in Subsection 1.1 with an inspec-  
313 tion and execution phase. The actual decision of which transition to be taken next is  
314 decided by the given SelectionStrategy instance to the execute method. This decision  
315 is delegated to the two methods setChoices and resolveChoice.

### 316 Dynamics of the ConcreteCmlBehavior

317 As mentioned multiple times the ConcreteCmlBehavior class is the default realization  
318 of the CmlBehavior interface and is the only one of them explained in details in this  
319 report. To understand the dynamic model we need to see what happens in the inspect






Figure 10: The top level dynamics

320 and execute methods, as these together determines the possible transitions at the top  
 321 level shown in the last section.

322 In Figure 11 the general inspect dynamics is depicted. When the inspect method is  
 323 called on a ConcreteCmlBehavior it uses its nextNode (in the java source nextNode  
 324 and nextContext is actually a Pair<INode, Context>) to delegate the actual inspection  
 325 to the CmlInspectionVisitor. The CmlInspectionVisitor contains a method case<INode  
 326 instance name> for every CML AST node. So e.g. if nextNode is a AInterleaving  
 327 gAction then the visitor method caseAInterleavingAction(..) method is called with the  
 nextContext. The called case method will return a Inspection which contains the next

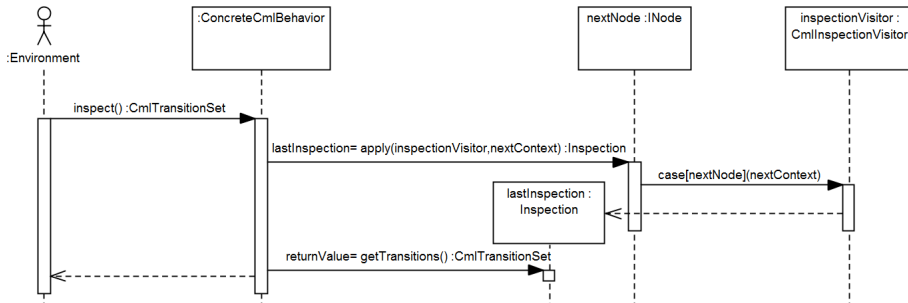


Figure 11: The general dynamics of the inspect method

possible transtions and a transition function to be called if the execute method is to be invoked. The last call in Figure 11 just grabs the CmlTransitionSet from the returned Inspection object and return this as the result of the inspect call.

The execute method, shown in Figure 12, will execute the given transition and must only be called if one the returned transitions from the inspect method has been chosen for execution. The actual execution is delegated to the CmlCalculationStep instance contained in the last calculated Inspection object (lastInspection in the figure) in the inspect method. The instance of a CmlCalculationStep is an anonymous class created

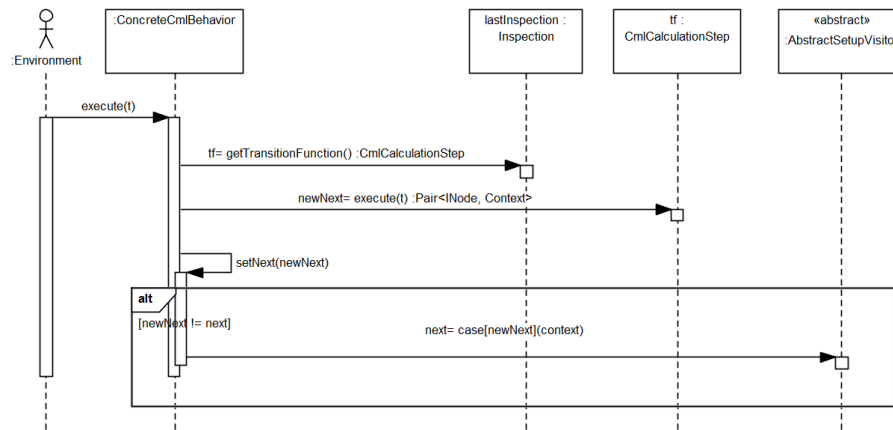


Figure 12: The general dynamics of the execute method

in a inspection visitor case, so the behavior of the execute method is entirely dependent on the current node contained in the next pair. The result of the execute method is the next node and context. As seen in Figure 12 the setup visitor is called if the newly returned pair is different from the current one. This enables any case specific setup behavior to be implemented here. E.g in some cases the context needs to be updated before the inspection phase is commenced.

### 3.3 The IDE Layer

This section will explain the IDE layer very briefly and only go through the very top level structure.

#### 3.3.1 Static Model

##### Packages

The following packages defines the top level structure of the IDE:

**eu.compassresearch.ide.interpreter.model** Contains all the classes that implements the Eclipse debug model [Wri04]

**eu.compassresearch.ide.interpreter.launching** Classes that deals with launching a debugging session.

353 **eu.compassresearch.ide.interpreter.protocol** Classes that deals with the communi-  
 354 cation between the Eclipse CML debugger and a CML interpreter instance.

355 **eu.compassresearch.ide.interpreter.view** Contains the custom views of the Eclipse  
 356 CML debugger.

357 Before explaining the steps involved in a debugging session, there are some important  
 358 classes worth mentioning:

359 **CmlDebugger** Interface with the responsibility of controlling the CmlInterpreter exe-  
 360 cution in a debugging session.

361 **SocketServerDebugger** Realization of the CmlDebugger interface that enables con-  
 362 trolling the debugging session over a tcp connection.

363 **DebugMain** Class that contains the main method that initializes the core component  
 364 on the debugger JVM side. This involves

365 **CmlDebugTarget** This class is part of the Eclipse debugging model. It has the re-  
 366 sponsibility of representing a running interpreter on the Eclipse side. All com-  
 367 munications to and from the Eclipse debugger are handled in this class.

### 368 3.3.2 Deployment Model

369 In order to get the big picture of how the IDE layer works together with the Core, a  
 370 deployment view of the IDE is shown in Figure 13.

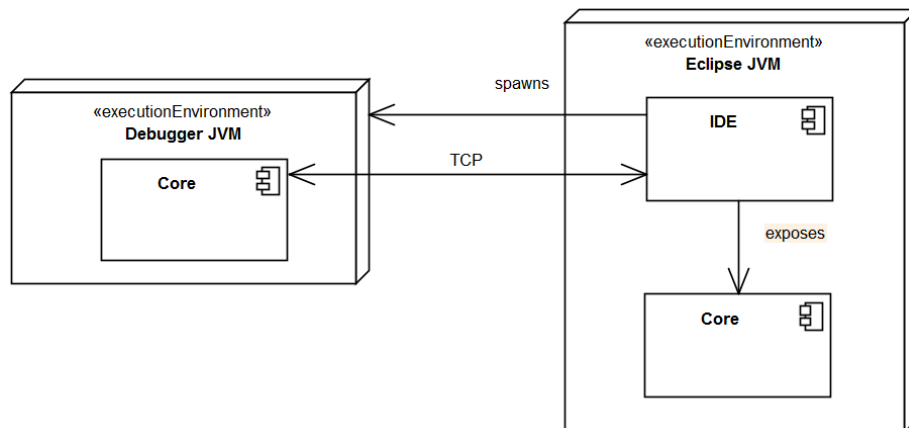


Figure 13: Deployment diagram of the debugger

371 An Eclipse debugging session involves two JVM instances, the one that the Eclipse  
 372 platform is executing in and one where only the Core executes in. All communication  
 373 between them is done via JSON through a TCP connection.

The JSON protocol need to be defined

375 A debugging session has the following steps:

- 376 1. The user launches a debug session

- 377 2. On the Eclipse JVM a CmlDebugTarget instance is created, which listens for an  
378 incoming TCP connection.
- 379 3. A Debugger JVM is spawned with the main method in the DebugMain class as  
380 starting point.
- 381 4. A SocketServerDebugger instance is created and tries to connect to the created  
382 connection from step 2.
- 383 5. When the connection is established, the SocketServerDebugger will send a START-  
384 ING status message along with additional details
- 385 6. The CmlDebugTarget updates the GUI accordingly.
- 386 7. When the interpreter is running, status messages will be sent from SocketServerDe-  
387 bugger and commands and request messages are sent from CmlDebugTarget.
- 388 8. This continues until either the CML model successfully terminates or the user  
389 stops.

## References

- [BGW13] Jeremy Bryans, Andy Galloway, and Jim Woodcock. CML definition. Technical report, COMPASS Deliverable, D23.4, September 2013.
- [COM] COMPASS. The cml syntax. <https://github.com/compassresearch/cml-syntax>.
- [GHJV94] Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1994.
- [Wri04] Darin Wright. How to write an eclipse debugger. <http://www.eclipse.org/articles/Article-Debugger/how-to.html>, aug 2004.